ME3 Voice Duplex: A Tool-Using, Full-Duplex Voice Agent for Raspberry Pi

Crizepvill F. Dumalaog
Mapua University
Master of Engineering in Artificial Intelligence
Manila, Philippines
202521406

Abstract

ME3 is a Raspberry Pi voice-agent prototype that combines local speech processing, full-duplex audio, a provider-selectable remote language-model brain (Qwen by default, with optional GPT-5.6 Luna), and project-owned Model Context Protocol (MCP) services. The system is designed for an important interaction beyond question answering: a user can ask by voice for a web application or other coding task, and the model can call bounded file, search, command, and background-job tools to perform the work. No task-specific model training is required. Instead, the system integrates pretrained Silero VAD, Qwen2-Audio-7B-Instruct ASR, Qwen3-TTS-12Hz-1.7B-CustomVoice TTS, and a retained but inactive ElevenLabs adapter, together with an OpenAI-compatible Qwen/vLLM endpoint or the OpenAI GPT-5.6 Luna API. The active Qwen audio path uses protected vLLM listeners; the text ToolBrain and bounded tools remain shared within the final release. This paper describes the architecture, tool-call protocol, safety boundary, local verification, Raspberry Pi release workflow, and evidence boundary for the separate FULL-DUPLEX-BENCH v1.5 evaluation. The current implementation is an engineering release, not a claim of completed paper-grade overlap results. The current revision also specifies the Vocalis-Pi event contract: continuous VAD/transcript inputs, explicit turn states, interrupt controls, and an adapter seam for asynchronous MCP harness work. The workstation acceptance includes the existing local checks plus mocked request-shape and configuration tests for the Qwen audio profile. An earlier VPN session authenticated the DGX3 Qwen ASR/TTS services and completed an injected voice/tool turn. On the latest laptop-hotspot check, however, the Pi OpenVPN unit is active but has no tunnel interface and cannot reach the DGX3 ports. The current accepted release is deployed at `/home/dalmacio/Desktop/dandan`; its Pi voice loop and HTTP/browser-audio endpoints start, but provider-live Qwen ASR/TTS are not claimed for this network session. A prior operator-confirmed human turn and the 30-trial injected benchmark remain separate evidence. The paired Pi/laptop launchers provide DHCP-independent mDNS discovery and SSH forwarding for the browser UI, speaker stream, and OAuth callback.

Keywords

voice agents, full duplex, Raspberry Pi, tool calling, MCP, Qwen, GPT-5.6 Luna, speech interaction, edge computing

1 Introduction

Voice interfaces are often presented as a sequence: wait for the user to stop, transcribe one utterance, generate text, and play a response. A useful assistant must also manage interruption and action. A user should be able to say "build a web app" and receive evidence of files and verification, rather than a plausible paragraph describing code that was never created.

ME3 addresses these requirements with a local audio loop and a remote tool-using language model. The design separates four concerns. First, the Pi keeps voice-sensitive processing local through VAD, ASR, and TTS. Second, a Qwen model served by an OpenAI-compatible vLLM API performs dialogue reasoning by default; GPT-5.6 Luna is an optional OpenAI-compatible alternative. Third, typed tools provide an auditable action surface. Fourth, a release and deployment workflow distinguishes local preparation from actual Pi installation.

The project is intentionally not a new model-training project. The agent behavior comes from inference and orchestration: pretrained speech models, Qwen or GPT-5.6 Luna inference, MCP/tool dispatch, and a bounded local execution policy. This choice reduces iteration time and makes the engineering questions explicit.

The contributions of this work are:

- a full-duplex local voice loop with barge-in during TTS playback;
- a provider-neutral OpenAI-compatible tool loop that can create and verify a static web application through voice or text;
- one bounded registry shared by direct Qwen calls and MCP wrappers;
- a user-intent-gated Telegram text adapter and Android call bridge; carrier SMS and Telegram voice calls are not implemented;
- background task supervision so builds do not block the microphone;
- a reproducible local-dev, package, deploy, and evidence workflow for the Raspberry Pi target, including a LAN browser-speaker transport and a vendored Full-Duplex-Bench v1/v1.5 adapter.

2 System Requirements and Scope

The required runtime contains four owned MCP services: VAD, ASR, TTS, and tasks. Pipecat is retained as the harness-facing integration path, while the local loop is the reference

implementation used for Windows development and the Pi service. The runtime configuration is kept in one YAML file so the audio rate, speech thresholds, model choices, endpoint, tool budget, and deployment path are not duplicated across scripts.

The current engineering scope includes interruptible spoken conversation, remote Qwen inference, project-local coding tools, public web lookup, persistent Pi-local timers, alarms and reminders, system information, the existing ME2 bridge, and credential-gated Telegram text and optional email actions. Google Calendar synchronization and Spotify playback are implemented but require account authorization; Spotify playback also requires Premium. Carrier SMS and Telegram voice calls are not available. The Android companion provides an unverified SIM-call bridge. ME3 does not claim unrestricted Onit equivalence. An always-listening device needs a narrower default action surface than an interactive coding IDE. Capabilities that would add external side effects, credential use, repository rewrites, or arbitrary shell access must be added as reviewed typed tools.

3 Architecture

3.1 Audio and conversation path

The default audio path is:

`microphone -> Silero VAD -> ASR -> Qwen/Luna -> fixed profile TTS -> speaker.`

Capture remains active while TTS plays. The local state machine has four states: `idle`, `listening`, `thinking`, and `speaking`. Speech during `speaking` stops playback and starts a new turn. A short pre-roll is retained so the beginning of an utterance is not clipped. The VAD endpoint requires a minimum speech interval followed by configured trailing silence.

The active Qwen profile always uses Qwen3-TTS CustomVoice `aiden` with a formal style instruction. Replies, backchannels, and completion announcements use the same profile voice. The Philippines-market profile keeps English, Filipino, and Taglish in the supported interaction set. The spoken policy is formal and professional: the fixed long-task acknowledgement is "I am processing your request now," and the assistant is instructed to avoid flirtatious, intimate, suggestive, teasing, or overly familiar language. The ElevenLabs adapter is retained as a documented future experiment but is not auto-selected by the active demo runtime. Qwen remains the default reasoning and tool-calling model; GPT-5.6 Luna is an optional provider using the same tool loop. ElevenLabs Realtime v1 Mini/Max is a separate hosted voice-agent architecture that also manages the LLM and ASR settings, so it is not silently substituted into this experiment.

On 26 October 2026, an authorized local provider-only test used the gitignored Windows development credential and an operator-selected account voice. ElevenLabs returned 54 raw-PCM chunks (163,468 bytes at 16 kHz), and a Filipino-language test phrase played through the workstation's default speaker. The adapter was corrected to reassemble 16-bit samples split across HTTP chunk boundaries and to classify chunked-encoding failures using the correct Requests exception handling. The account voices inspected for this test were English-labelled; therefore this result verifies the provider integration and output path, not approval of a native Filipino accent.

The duplex state machine is independent of the Qwen model. This makes the interaction control testable without a network request and leaves room for a Pipecat transport to reuse the same transition logic.

For the ME2-style Pi web experience, capture remains on the Pi while an optional browser-output transport publishes generated 16 kHz mono PCM over a LAN WebSocket. The browser schedules that stream through Web Audio and can select the PC's speaker; alternatively, the same UI selects a Pi-local output device. The HTTP dashboard exposes safe PortAudio device names and restarts only the audio loop when a route changes. A one-click browser connection is required by autoplay policy, but it does not create a push-to-talk interaction: the Pi microphone remains continuously active and barge-in is handled by the Pi state machine.

The selectable `qwen_instruct` profile changes only the ASR/TTS providers. Local Silero VAD and the same duplex state machine accept barge-in. Completed turns are sent to `Qwen/Qwen2-Audio-7B-Instruct` through an OpenAI-compatible multimodal `/v1/chat/completions` request carrying the WAV as an `audio_url` data URL. Replies are rendered by `Qwen/Qwen3-TTS-12Hz-1.7B-CustomVoice` through `/v1/audio/speech`. The text ToolBrain remains the existing `Qwen/Qwen3.8-27B` or optional Luna provider, so requests such as "build a web app" retain the same file and verification tools. Endpoint and key values are runtime-only. An ASR/TTS pair is therefore not a replacement for the general text tool-calling model.

3.2 Vocalis-Pi event contract

The implementation makes the full-duplex behavior observable through a transport-independent Vocalis-Pi orchestrator. Its external states are `LISTENING`, `THINKING`, `SPEAKING`, `BACKCHANNELING`, and `INTERRUPTED`. The local loop remains active on the microphone while the speaker plays. A VAD update received during `SPEAKING`, `THINKING`, or `BACKCHANNELING` publishes a `barge_in_event`, then emits `FLUSH_QUEUE` and `FADE_OUT` TTS controls before returning to `LISTENING`.

The protocol carries four input/event types: VAD state, transcript deltas, overlap signals, and asynchronous harness events. Its output types are assistant-text deltas, TTS control actions, turn-state transitions, and optional harness dispatch requests. A long Qwen tool turn can therefore produce a short `BACKCHANNELING` acknowledgement while its task continues. Tool progress is mirrored as sanitized harness events; arguments and results are excluded from these events because they can contain user data or credentials. The exact channel names and JSON examples are documented in the repository's Vocalis protocol manual.

The optional harness adapter emits a JSON-RPC 2.0 `tools/call` request named `harness_dispatch`. The default

verified release keeps the bounded Qwen tool loop in-process for low overhead and uses the adapter as a stable seam for an Onit, Pipecat, or OpenCode background worker. This makes the contract testable without claiming that a separate external harness is already running in every deployment.

3.3 Provider-neutral tool-using brain

The remote brain is an OpenAI-compatible chat-completions client. The first request contains the recent dialogue and JSON Schema definitions for the ME3 tools. Qwen or GPT-5.6 Luna may return assistant text or a list of function calls. For every call, ME3 parses the arguments, executes the local function, appends a `role=tool` message, and asks Qwen to continue. The checked-in configuration allows 24 tool rounds for demanding multi-file or web-app tasks. After the final permitted batch, ME3 makes one no-tool summarization request instead of exposing an internal tool-limit error to the user.

For Qwen, the checked-in model is served by vLLM. For Luna, the model ID is `gpt-5.6-luna`; hidden reasoning is disabled, temperature remains omitted, and the client requests a direct response stream when supported. The protocol is illustrated below.

```
{
  "model": "Qwen/Qwen3.8-27B",
  "messages": [{"role":"user", "content":"build a web app"}],
  "tools": [{"type":"function", "function":
    {"name":"write_file", "parameters": {}}}],
  "tool_choice": "auto"
}
```

After a tool call, the result is returned as structured JSON. Qwen therefore receives the actual file path, byte count, return code, or error. The final spoken response is not accepted as evidence by itself.

3.4 Owned MCP services

The VAD server exposes frame and stream detection plus endpointing. The ASR server exposes PCM and file transcription. The TTS server exposes synthesis, streaming chunks, and installed voices. The tasks server retains timers, system information, and the ME2 bridge, and now wraps the same bounded project tool registry used by the direct Qwen path. The shared implementation avoids one policy for the local loop and another policy for MCP clients.

4 Tool Registry and Safety Boundary

Table 1 summarizes the direct functions visible to Qwen.

The path policy resolves every path inside the ME3 project. Credential-like names, environment files, private keys, the Git directory, the virtual environment, model caches, and generated release trees are protected. Reads and writes are size bounded. Commands are tokenized and run without a shell; shell operators and patterns associated with deletion, privilege escalation, SSH/SCP, repository reset/clean/push, and package publication are rejected. Long commands run with standard input closed, captured output, and a timeout.

Table 1: Direct Qwen tools in the ME3 registry

Tool	Purpose
`list_files`	List project files and sizes.
`read_file`	Read bounded project text.
`write_file`	Create or replace a project text file.
`search_text`	Search project text files.
`run_command`	Run a direct allowlisted command.
`start_background_task`	Start a bounded build or test.
`task_status/output/cancel`	Supervise a background task.
`system_info`	Return non-secret host information.
`fetch_url/web_search`	Read public HTTP(S) resources.
`git_status`	Inspect repository status without mutation.
`channel_health`	Report safe Telegram and optional email configuration.
`send_telegram_message/send_email`	Confirmation-gated external communication through Telegram Bot API and optional Resend.

This boundary is a design choice rather than an implementation accident. A voice transcript can be ambiguous, and the assistant is not an appropriate place for silent arbitrary shell execution. The registry is sufficient for the demonstration web app while keeping higher-risk operations explicit.

The external communication tools use the same registry and are stricter than the local project tools. `channel_health` is network-free. Each `send_telegram_message` or optional `send_email` request first validates its target and returns a confirmation-required status unless the operator has explicitly confirmed the side effect. The adapters then send one bounded REST request and return only a status, channel, and safe provider ID; provider response bodies and authorization headers are discarded. Ordinary Telegram voice calls are not available through the Bot API. The Android companion can receive an explicit call request over the browser-audio WebSocket and invoke the phone's dialer/call permission; a live call has not yet been verified. Cellular SMS is not implemented. Inbound webhooks, identity/CRM actions, consent, recording, and human transfer remain outside this adapter boundary.

5 Web-App Capability

The target workflow is a dependency-free web application when the user does not name a framework. Qwen writes an application directory under `agent_workspace/<name>`, lists the result, reads the important files, and runs a bounded verification command if one is appropriate. The provided demonstration contains `index.html`, `styles.css`, and `app.js`; it

provides a command input, microphone placeholder, status area, and tool-activity log.

The capability is demonstrated twice: the first live request used Qwen to create the HTML file through a real tool call; a later live request used Qwen to list the directory and read all three files. This demonstrates inference, tool dispatch, file mutation, and evidence retrieval without training a model. The demo is intentionally local and static; the current safety policy does not publish it to the public internet or install an arbitrary server.

6 Deployment Workflow

ME3 follows a four-stage local-dev workflow:

(1) **Environment gate:** import the shared Python dependencies.
(2) **Smoke gate:** load all four MCP services, call their health tools, construct the engineering benchmark, and query PortAudio when available. The vendored public fd-benchmark is also available through a direct-feed ME3 adapter; its simulation is used as a protocol smoke test, not as a hardware result.
(3) **Package:** copy the runtime files into a versioned release, write SHA-256 entries, and create a zip archive.
(4) **Deploy:** transfer the archive to the configured Pi, merge its stable `me3_voice_duplex` directory, run the idempotent installer, and record a deployment receipt.

The current target is dalmacio@cfdfnjrpi5.local at /home/dalmacio/Desktop/dandan. Its Wi-Fi address is DHCP-assigned; the hostname is the stable same-LAN route. The retired `dandan-qwen` directory was moved to a dated hidden archive, leaving one active ME3 service. The current Pi release reports `loop:true`; tunneled HTTP health and a browser-audio WebSocket handshake pass, and the web clock API creates/cancels a timer. The SHARC VPN service is enabled but its current hotspot route cannot reach DGX3. These are deployment and interface checks, not a fresh human spoken turn or a provider-live Qwen claim for the current network.

The API key is never included in the archive. On the Pi it must be provisioned through a secure runtime environment. The Windows DPAPI helper can report configured/not-configured and can provide an in-memory child-process value for local testing, but its secret is not printed or written to the repository.

7 Verification and Evidence

The local verification gate runs the shared-environment check, audio/MCP smoke test, paper-alignment check, unit suite, and a live Qwen tool turn when its provider endpoint is reachable.

The active Qwen test suite passed 99 tests in the current accepted checkout. The tool-focused tests cover OpenAI-compatible schemas, blocked shell patterns, background task completion, and a mocked multi-step Qwen response. The ElevenLabs tests additionally cover split-byte PCM reassembly, genuinely odd total PCM rejection, timeout retry, and chunked-encoding retry without credential disclosure. The channel tests cover confirmation gating, Telegram chat-id validation, Bot API request construction, Resend request construction, and safe metadata. Python compilation also passed. The Vocalis tests cover state transitions, JSON-RPC harness dispatch, barge-in TTS-control ordering, and LocalVoiceLoop tool-progress bridging. One current injection-only fixture run also completed with one turn and zero errors; it remains engineering evidence and not physical-speaker or v1.5 evidence. The live Qwen smoke test records an actual tool call and final answer. The local smoke gate loads the four MCP servers and confirms the current Windows PortAudio inventory. The integrated fd-benchmark simulation also completed a 10-trial development-dataset run, yielding a protocol-and-report smoke result with no ASR judge; this is not a paper-grade or physical Pi result. These are engineering checks, not a claim about Pi audio.

The evidence status at this revision is summarized in Table 2.

The deployed Pi preflight is also an explicit hardware gate. It is read-only, does not alter mixer state, and returns exit code 2 when no capture device is visible. This keeps a playback-only Pi from being mislabeled as a working full-duplex voice endpoint.

When speaker hardware is unavailable, the project uses a separate injection-only lane. The evaluator feeds recorded 16 kHz PCM into the same VAD, ASR, Qwen/tool, TTS, and interruption state machine at wall-clock frame cadence, while a virtual sink records TTS timing without opening PortAudio output. This verifies injected-input agent behavior and virtual barge-in, but does not claim physical-speaker behavior or paper-native v1.5 evidence.

To align the class evaluation with the supplied public benchmark, the release vendors the MIT-licensed fd-benchmark repository and adds `scripts/me3_fdbench_agent.py`. The adapter connects the real `LocalVoiceLoop` to the benchmark's direct-feed WebSocket: incoming 16 kHz mono int16 frames enter the same VAD/ASR/tool/TTS path, outgoing PCM is recorded by the benchmark rig, and a barge-in emits its queue-clear event. A small simulation run verifies the integration without opening a microphone or speaker; full datasets, ASR/judge scoring, and a clean Pi run remain required for publication-quality measurements.

The project also contains a separate engineering latency driver and a paper-native FULL-DUPLEX-BENCH v1.5 contract. The v1.5 contract includes four scenarios with 200, 99, 100, and 100 stimuli and response categories RESPOND, RESUME, UNCERTAIN, and UNKNOWN. The engineering five-scenario driver must not be presented as the paper's 499-utterance result. A complete captured overlap artifact, transcript classification, and prosodic before/after analysis are still required before claiming a paper-grade result.

Table 2: ME3 progress evidence at the 7 November 2026 revision

Evidence lane	Status and interpretation
Tool-capable Qwen	Earlier live turns created and inspected a web app through bounded tools.
Audio and deployment	Pi loop, SSH-tunneled health, and browser-audio WebSocket pass; no fresh human turn on this release. Windows loopback barge-in was tested earlier.
Class latency stage	30 trials were retained: 29 replies plus one no-reply; the aggregate used faster-whisper fallback, and barge-in was not measured.
Qwen audio services	Earlier VPN probes measured Qwen ASR 431.9 ms and TTS 851.6 ms; the current hotspot cannot reach DGX3.
Clock and channels	Schedule create/cancel passes; no alarm was fired. Live Pi reports Spotify, Calendar, and Telegram unconfigured; SIM calls are untested and carrier SMS is absent.
Paper-native FDB v1.5	No clean 499-utterance overlap artifact or paper-grade result is claimed.

8 Discussion

The main engineering trade-off is locality versus model capability. Keeping speech I/O local reduces the need to upload audio and makes the Pi responsive, while Qwen supplies stronger reasoning and code generation than a small local model. The cost is network dependence, endpoint latency, and the need to provision the API key on the target device. The typed registry makes the boundary visible and testable, but it is not a complete replacement for an interactive coding environment.

The local ElevenLabs result narrows one implementation risk: the selected streaming endpoint can produce audio in the PCM format consumed by the full-duplex playback path, and the adapter tolerates arbitrary HTTP chunk boundaries. It is not part of the active single-voice runtime. A complete voice-agent turn still requires both a reachable Qwen brain and an operator approved voice profile for the intended Filipino customer-support deployment. The new communication adapters narrow the integration boundary but do not establish that the agent can safely replace a staffed call center: provider accounts, webhooks, media streaming, consent, retention, escalation, and supervised quality gates remain necessary. Provider rendering, Qwen/tool execution, and external-channel side effects are therefore reported as separate evidence lanes rather than collapsed into one end-to-end claim.

The full-duplex behavior is also more than a transport setting. VAD must keep running during playback, the audio output must be cancellable, and the turn-state transition must preserve the interrupted user's utterance. Echo control, microphone placement, and ALSA device selection remain physical Pi concerns. A successful Windows loopback does not prove those conditions.

Future work should add a Pi-local live self-test, Telegram inbound webhook signature verification, secure credential provisioning and rotation, an operator-owned Telegram/Resend test, a reviewed web-app preview server with explicit approval, automatic physical-speaker barge-in/echo evaluation, and the full v1.5 overlap capture/evaluation. Additional Onit-like capabilities should be introduced as typed, auditable functions rather than by removing the voice safety boundary.

9 Conclusion

ME3 now combines full-duplex voice interaction with a real Qwen tool loop and an explicit Vocalis-Pi event contract. A spoken build request can result in actual project files and verification, and the agent can supervise longer work without freezing the microphone. The implementation uses inference and tools rather than model training, keeps local speech services under project ownership, and records the distinction between local verification, prepared release, Pi deployment, and benchmark evidence. The ElevenLabs adapter was tested locally but is not part of the active single-voice runtime. The latest workstation acceptance passes 99 tests and the class injection benchmark reports aggregate mean and p95 over 29 replies plus one retained no-reply trial. The Pi browser-speaker path, schedule API, and vendored benchmark adapter pass current protocol checks. Qwen ASR/TTS were provider-live during an earlier VPN session but are unreachable from the current hotspot. Remaining decisive checks are restoring DGX3 access, a fresh human Pi spoken turn, automatic physical-speaker barge-in, live channel setup, and a clean captured v1.5 overlap evaluation.

References

[1] Lin et al. "FULL-DUPLEX-BENCH v1.5: Evaluating Overlap." Project source PDF: *FULL-DUPLEX-BENCH V1.5 - EVALUATING OVERLAP.pdf*, AI 231 ME3 repository, 2026.

[2] "Building AI Voice Agents - Summary." Project source PDF, AI 231 ME3 repository, 2026.

[3] A. Radford et al. 2023. Robust speech recognition via large-scale weak supervision. In *Proceedings of ICML*.

[4] Silero Team. Silero VAD. Open-source voice activity detection model and runtime. The ME3 release pins the installed runtime version in `requirements.txt`.

[5] Model Context Protocol. Official protocol specification and SDK. The ME3 release uses the Python MCP/FastMCP stack for owned task services.

[6] Pipecat. Open-source framework for real-time voice and multimodal agents. The ME3 release pins Pipecat 1.13.0 for the harness-facing adapter.

[7] ElevenLabs. "Realtime models" and "Stream speech." Official developer documentation, 2026. ElevenLabs developer documentation.

[8] Qwen Team. "Qwen2-Audio-7B-Instruct." Official model card, 2024. Hugging Face model card.

[9] vLLM Project. "Qwen3-TTS online serving." Official vLLM-Omni recipe, 2026. vLLM-Omni recipe.

[10] Resend. "Send Email." Official API reference, 2026. Resend API reference.